

Report on Ping Pong Time Test

Project 1 part 1 and 2

February 20, 2017

Evan Su

Question:

Which method of sending message would be faster, message passing between processes in the same box or message passing between processes in different boxes?

Test Methodology:

A program, called `mpi_pingpong`, is written to test the send and receive time between processes. The program will measure the time elapsed for message sent to another process and back to the sender. The size of the message is 2 integers. The size of 2 integers was arbitrary chosen. The program will measure the time of 50 sends and receives. The results are outputted on the terminal in the format of CSV (comma separated values). For each trial, the program will run twice; once for the same box and another for different boxes.

Data Analysis:

The results of the test are summarized in Table 1.1. Outliers from the results were removed. The results show that message passing in the same box is about 2 - 3 microseconds. Whereas, message passing between different boxes is around 52 microseconds. The fast speed of same box message passing is a result of shared memory. Since message passing occurred in the same box, the processes did not have to incur network traffic.

Conclusion:

Message passing between processes in the same box is significantly faster than in different boxes. The same box is at least 20 times faster. The high time cost of network usage implies that any message passing should ideally be between processes in the same box.

Graphs and Tables

Table 1.1

Trial	One Box					Two Box			
	MIN	MAX	Average	STDEV		MIN	MAX	Average	STDEV
1	1	5	2.142857	0.645497		57	73	59.71429	3.5
2	2	3	2.291667	0.45934		46	60	52.16327	3.52578
3	2	9	2.857143	1.040833		49	81	52.91837	6.214354

Time is measured in microseconds

Report on Packet Size Test

Project 1 part 3

February 20, 2017

Evan Su

Question:

How large does the message need to be for message passing routine to split the message into multiple packets?

Test Methodology:

A program, called `mpi_packetCheck`, is written to test time elapsed for a message sent to another process and back to the sender. The program will test each possible message size up to 10,000 integers. Message size of 10,000 integers was chosen because the data from 1,000 integers was insufficient and 100,000 integers used too much computational time. The program sends integers between two process on two different boxes to measure the communication time. Each message size test ran 50 times before the program tests the next amount.

Data Analysis:

Four trials were performed and the results were plotted on graph 2.1. Each trial ran on different days on different times. Trial 1 had the least network activity of 1 user. Trail 4 had the most network activity of 5 users.

The communication time linearly increased until 500 integers. The message passing routines was creating packets of exact message sizes. The communication time then increased linearly but at a slower rate until 1947 integers. The message passing routines increased the size of the packets linearly but not exact sized anymore. After 1947 integers, the communication time continually spiked up, plateau. At this point, messages began to be split into different packets.

The trials had time spikes at different places but the size of the spikes was consistent. Each time spike is 143 microseconds in height. This implied that each time a spike occurred, an additional packet was made and sent. The number of packets can be determined by diving the time for a message by 143 microseconds. For example, between 2000 to 2780 integers, all trials used at least 2 packets. Looking at the upper bound of each number of packet interval for all trials (1800, 2800, 5005, etc.), all packet carried at most 2^{16} bits of data. For instance, around 2800 integers, the number of packets that carry the message jumped from 2 packets to 3 packets. Before the jump, the 2 packets must be full which implies that each packet is carry 1400 integers. 1400 integers are approximately $2^{15.4}$ bits. Testing each upper bound show that no packet will exceed 2^{16} bits.

Conclusion:

The maximum packet size appears to be 2^{16} bits. The send routine increases packet size linearly until the limit is reached. The send routine will always split the message into another packet before reaching the limit.

Graphs and Tables

Graph 2.1

